

BRICK 2.0 → BRICK-AM

A walkthrough of the updates — data sources, justifications, and code changes for the Antarctic-Mengel recalibration

M. Sarofim · NYU Marron Institute — compiled 2026-07-24 · repo `SLR-RFF-BRICK`, branch `brick-mengel-vnext` (HEAD `69a2bfb`)

0. Orientation

BRICK-AM (“Antarctic-Mengel”) is our fully-updated MimiBRICK sea-level model. It differs from **BRICK 2.0** (Tony Wong’s obs-driven MimiBRICK v2.0.0 port) in three substantive ways plus one numerical correction:

#	Update	In BRICK 2.0?	In v2.1.0?	Category
1	Mengel-2016 glacier emulator (replaces the Wigley–Raper single-reservoir glacier)	No	Yes	model structure
2	Observation updates (Dangendorf/STAR total; GRACE-FO + GlaMBIE extensions; proper σ ; dropped point terms)	No	No — v2.1.0 keeps Church–White	calibration data
3	GMST→Antarctic amplification a : equilibrium 1.196 → CMIP6 secant 1.08 ± 0.15	No	No	calibration prior
—	Sub-annual DAIS crossing correction (a numerical fix, only affects the pulse marginal)	No	No	numerics

The name emphasizes update 3: a decomposition (§7) attributes **~85 % of the level change** from BRICK 2.0 to the amplification, with the glacier and observation updates far smaller. Relative to Tony Wong’s official **v2.1.0** (which already carries the Mengel glacier), the BRICK-AM changes are **updates 2 and 3** — see below.

BRICK 2.1 and the baseline

Tony Wong’s official MimiBRICK releases (confirmed by Tony) are: **v2.0.0 = Wigley–Raper glaciers + Church–White GMSL**; **v2.1.0 = Mengel glaciers + Church–White GMSL**. So the Mengel glacier (**update 1**) is **shared with the official v2.1.0**, and relative to v2.1.0 the BRICK-AM changes are **updates 2 and 3** (plus the numerical fix). A large part of update 2 is precisely the total-GMSL observation swap **Church–**

White → **Dangendorf 2024** + **NOAA STAR** (§3), alongside the GRACE-FO / GlaMBIE component extensions and the Frederikse-ensemble σ .

This document is written from the **BRICK 2.0** baseline (all three updates visible) because that is where our development actually happened and what the decomposition (§7) measures. Two implementation facts to keep straight:

- Our BRICK-AM is built on **stock MimiBRICK v2.0.0 + a local Mengel component** (`julia/glaciers_mengel_component.jl`) swapped in at runtime (§3a) — we did not build on the packaged v2.1.0. The Mengel *structure* matches v2.1.0; our *posterior* differs, because we calibrate the same emulator against different targets (the extended Dangendorf / GRACE-FO / GlaMBIE set, not Church–White).
- The `v2.1` git tag in SLR–RFF–BRICK is **our own repo release tag**, unrelated to the MimiBRICK package version. The Julia depot pins `1.0.1`, the `v1.2.1` tag, and the `v2.0.0` tag (slot `edplP`); the BRICK-AM calibration runs under the `julia_v2` environment (= v2.0.0). The Mengel model is also canonicalized in the separate **MimiBRICK-FM** repo.

BRICK version map

Repo / branch or env	MimiBRICK build	Role
SLR–RFF–BRICK branch <code>brick-v1.2-vehicle</code>	v1.2.1 (<code>RtLCv</code>)	pre-2.0, EPA-rescission-comparable arm
SLR–RFF–BRICK branch <code>main</code>	v2.0.0 (<code>edplP</code>)	BRICK 2.0 = Wong obs-driven port
SLR–RFF–BRICK branch <code>brick-mengel</code> (archived, tag <code>archive/brick-mengel-2026-06-17</code>)	v2.0.0	frozen prior Mengel/study-driver work
SLR–RFF–BRICK branch <code>brick-mengel-vnext</code> (env <code>julia_v2</code>)	v2.0.0 + local Mengel	BRICK-AM (this document)

Reproducibility flag. `brick-mengel-vnext` is **local-only** (not pushed to origin). Anyone reproducing this needs the working tree, not a clone of origin.

1. The three updates at a glance

Each update is developed below in its own section with four parts: **what changed**, **data sources** (with paths + provenance), **justification** (with the authoritative write-up), and **code** (with paths). File paths are repo-relative to `SLR–RFF–BRICK/` unless prefixed `FaIRtoFrEDI/`.

The **forcing** is common to all rungs: **FaIR 2.2.4 (calib1.4.5)** ensemble-mean GMST + OHC on Smith-harmonized SSP2-4.5, files `data/observations/fair_mean_gmst_ssp245harm.csv` and `..._ohc_ssp245harm.csv`, built in the sibling repo:

```
# in FaIRtoFrEDI/
python build_fair_mean_v145.py \
  --emissions-file calibration_v145/emissions_v145_ssp245_harmonized.csv \
  --tag ssp245harm --scenario-label ssp245_harmonized
```

The harmonized splice (not the RCMIP-native `fair_mean_*_ssp245.csv`) is used deliberately so the calibration fit window and the pulse projections sit on identical forcing (`calibrate_mcmc_ext.jl:85–95`, `FORCING_TAG="ssp245harm"`).

2. Update 1 — The Mengel-2016 glacier emulator

What changed. BRICK 2.0’s glacier & small-ice-cap (GSIC) component is the Wigley–Raper single-reservoir model. BRICK-AM replaces it with the **Mengel et al. 2016** (PNAS 113:2597) two-timescale emulator, driven by total GMST, with a Little-Ice-Age (LIA) disequilibrium baseline.

Data sources

The Mengel emulator ingests **no external glacier time series** — it is forced by GMST and models the post-LIA committed melt internally (see below). The only glacier *observation* is the calibration **target** (§3, update 2): the `gsic` column of `outputs/recalib_targets_ext.csv`, a splice of **Frederikse 2020** (1900–2018) and **GlaMBIE 2025** (2000→2023; DOI 10.5904/wgms-glambie-2024-07, raw `data/observations/raw/glambie_global_glacier_mass.csv`).

Justification

The Wigley–Raper glacier structurally undershoots the 20th-century record (the “GSIC H1/H2 structural undershoot”). Mengel’s design fixes this by construction: with `gic_T_lia < 0` the equilibrium contribution `S_eq(T=0) > 0`, so glaciers are out of equilibrium at the 1850–1900 baseline and are *committed* to post-LIA melt — reproducing the early-20th-century melt directly, without an anthropogenic/natural forcing split or an external “natural” budget. A prototype validated the emulator against the Frederikse glacier total to $RMSE \approx 0.13$ cm. The full rationale and prototype fits are in `notes/handoff_2026-06-12_brick_mengel_calibration.md` and `handoff_2026-06-13_brick_mengel_post2018_extension.md`.

Code

File	Role
julia/glaciers_mengel_component.jl	the <code>@defcomp glaciers_mengel</code> (equations below); output var <code>gsic_sea_level</code> matches the WRB component so wiring into <code>global_sea_level</code> is unchanged
julia/brick_mengel.jl	<code>build_brick_mengel(...)</code> builds a stock v2.0.0 model then <code>replace!(m, :glaciers_small_icecaps => glaciers_mengel) (:51–52)</code> ; <code>update_brick_mengel!</code> sets the 7 Mengel params + non-glacier params
julia/brick_param_updates.jl	<code>update_brick_params!(...; skip_glaciers=true)</code> skips the WRB glacier mapping when Mengel is in place (:79–84)
python/calibrate_mengel_glacier.py, python/glacier_2tau_validate.py	offline prototype least-squares fits (single- τ and 2- τ) that informed the prior <i>ranges</i>

Equations (`glaciers_mengel_component.jl`:56–59), with T = total GMST anomaly rel. 1850–1900:

```
S_eq[t] = gic_a * (1 - exp(-gic_b * (T[t-1] - gic_T_lia)))
S_fast[t] = S_fast[t-1] + (gic_f * S_eq[t] - S_fast[t-1]) / gic_tau_fast
S_slow[t] = S_slow[t-1] + ((1-gic_f) * S_eq[t] - S_slow[t-1]) / gic_tau_slow
gsic_sea_level[t] = S_fast[t] + S_slow[t]
```

Parameters (:37–44): `gic_a` (asymptotic max SLE from the LIA state, m), `gic_b` (T-sensitivity, 1/K), `gic_T_lia` (LIA equilibrium temp, °C < 0), `gic_f` (fast-mode fraction), `gic_tau_fast` / `gic_tau_slow` (yr), `gic_sl0` (init m).

Important: the six Mengel parameters are **freed and sampled in the MCMC** (§6), not fixed. Their “central” values are the MCMC **prior means** (`calibrate_mcmc_ext.jl`:146–151: `gic_a` N(0.45,0.08); `gic_b` N(0.52,0.25); `gic_T_lia` N(−0.45,0.30); `gic_f` N(0.50,0.30); `gic_tau_fast` N(40,30); `gic_tau_slow` N(300,200)) and the medoid init (`a=0.45, b=0.52, T_lia=−0.45, f=0.5,  $\tau_{fast}=40, \tau_{slow}=250$`). The offline python fit `outputs/mengel_glacier_2tau_params.csv` **railed at multiple bounds** and was used only to set prior ranges — do not treat it as the deployed parameterization. (That offline railing is a *separate* thing from the MCMC posterior below.)

Posterior constraint (extA108). The deployed posterior is genuinely informed by the 1900–2023 `gsic` series, **not** prior-dominated. Every glacier parameter tightens **1.5–6× over its prior** (posterior/prior σ -ratio 0.16–0.69) and every posterior mean moves off the prior — e.g. `gic_b` 0.52→0.86, `gic_T_lia` −0.45→−0.95, `gic_tau_fast` 40→58. The committed-melt combination that actually drives the contribution, $S_{eq}(0)=a \cdot (1-\exp(b \cdot T_{lia}))$, is pinned to **0.20 ± 0.02 m SLE (12 %)**. **One prior was too tight:** `gic_T_lia` rails against its −1.0 °C lower bound (29 % of the posterior within 2 % of it, 5th percentile *at* the bound) — the data prefers a colder LIA equilibrium than the prior allows (`gic_a` / `gic_b` graze their bounds mildly, ~8–9 %). The posterior correlations are weak (strongest `gic_a` ↔ `gic_f` =

−0.61), so this is a genuine pull toward colder, not a degenerate ridge sliding into a corner. **Candidate follow-up:** widen the `gic_T_lia` lower bound (−1.5/−2.0) and re-check whether the posterior moves interior and the fit improves — the well-constrained `S_eq(0)` combination is unaffected regardless.

3. Update 2 — Observation updates

What changed. The calibration targets were rebuilt into a **reconciled, extended-to-present** product; the uncertainty band was put on a defensible footing (Frederikse 5000-member ensemble); and the legacy IMBIE/Dyurgerov Gaussian *point* terms were dropped in favor of the extended time series constraining the modern rate directly.

Data sources — `outputs/recalib_targets_ext.csv` (+ `_sources.csv`)

Built by `python/prep_recilib_targets_ext.py`. Series are 1900–2026, cm rel. the **1995–2005** window, NaN where a component lacks data. Splices are **offset-match only** (a level shift over the overlap window, no rescale; `GT_PER_CM_SLE = 3620.0`).

Component	Historical	Modern extension → end yr	Raw file / DOI
AIS	Frederikse 2020 (1900–2018)	GRACE-FO JPL mascon RL06.3Mv4 → 2026	<code>raw/grace_antarctica_mass.txt</code> · DOI 10.5067/TEMSC-3JC634
GIS	Frederikse 2020	GRACE-FO JPL mascon → 2026	<code>raw/grace_greenland_mass.txt</code> · same DOI
GSIC	Frederikse 2020	GlaMBIE 2025 → 2023	<code>raw/glambie_global_glacier_mass.csv</code> · DOI 10.5904/wgms-glambie-2024-07
Steric / TE	Frederikse 2020	NOAA NCEI 0–2000 m thermosteric → 2025	<code>raw/noaa_thermosteric_w0-2000m_yearly.dat</code>
Total	Dangendorf 2024 GMSL reconstruction (1900–2021)	NOAA STAR altimetry → 2024	<code>data/observations/dangendorf2024_gmsl_annual.csv</code> + <code>nasa_gmsl_annual.csv</code>

Two provenance points worth stating explicitly:

- **The “total” is genuinely Dangendorf-independent, not Frederikse.** The 2026-07-20 M3 rework (`prep_recilib_targets_ext.py:156–166`, commit `cda7ca2`) replaced an earlier total that was secretly the Frederikse sum with the real Dangendorf 2024 reconstruction. (This closed the “Dangendorf/Frederikse mislabel” issue.)

- **IMBIE 2023 (Otosaka et al.) is an independent cross-check only**, never fed to the fit (`prep_recalib_targets_ext.py:18–19, :258–271`).

Uncertainty σ — from the **Frederikse 2020 5000-member weighted component ensemble**

`data/observations/raw/frederikse2020_GMSL_ensembles.nc` (redistributed in Dangendorf’s Zenodo 10621070). `load_ensemble_sigma()` (`:94–113`) re-references each member to 1995–2005, takes the per-year weighted sd, and writes `value ± 1.645· $\sigma$` into `_lo/_hi` so the Julia likelihood recovers σ exactly via `eband=(hi-lo)/(2·1.645)` (`calibrate_mcmc_ext.jl:97`). The Dangendorf total borrows the ensemble GMSL sd because Dangendorf’s own per-year SE is corrupted upstream.

Rignot 2019 SMB anchor — enters not as a target column but as a likelihood term (§4, A5): grounded-AIS SMB 2098 ± 133 Gt/yr, area-scaled $\times(10.92/12.295) = 0.888 \rightarrow 1863.4 \pm 118.1$ Gt/yr (`calibrate_mcmc_ext.jl:230–245`).

Justification

Extending the AIS/GSIC series to the present lets the **time series** constrain the modern melt rate directly, which is both more informative and avoids double-weighting the same information via a separate Gaussian point term:

“DROPS the IMBIE dAIS(92-17) + Dyurgerov dGSIC(61-03) Gaussian point terms: the extended AIS/GSIC time-series now constrain the modern rate directly ... avoids double-weighting” — `calibrate_mcmc_ext.jl:22–24`

The phase-2 wiring (targets M3 + likelihood A2/A4/A5) is documented in `notes/handoff_2026–07–20_phase2_calibration_wired.md` and `prerun_summary_2026–07–20_phase2_calibration.md`.

Code

File	Role
<code>python/prep_recalib_targets_ext.py</code>	builds <code>recalib_targets_ext.csv</code> + <code>_sources.csv</code> from the raw obs; splices + ensemble σ
<code>julia/calibrate_mcmc_ext.jl</code>	the likelihood: AR(1) heteroscedastic <code>hetero_logl_ar1</code> (<code>:74–81</code>), per-series obs-end years (<code>:99–130</code>), Rignot anchor (<code>:290–292</code>), point terms absent from <code>logposterior</code> (<code>:264–303</code>)

The likelihood covariance for each of the five series is $\Sigma = \sigma^2/(1-\rho^2) \cdot \rho^{|i-j|} + \text{diag}(\epsilon^2)$ — an AR(1) process with per-year heteroscedastic observation error ϵ (the ensemble σ). Each series has its own free `sd_s`, `rho_s` (10 AR(1) nuisance params).

4. Update 3 — The GMST→Antarctic amplification

What changed. The DAIS component drives Antarctic ice-sheet temperature from global GMST through a linear map with slope `a` (the “amplification”). BRICK 2.0 uses the Shaffer-2014 **equilibrium/paleo** slope, `a`

$= 1/0.8365 = 1.19546$. BRICK-AM recalibrates it to the **CMIP6 transient secant**, prior $a \sim N(1.08, 0.15)$.

The mapping (and what the constants mean)

The DAIS component computes Antarctic surface temperature as

(`julia/patches/antarctic_icesheet_smoothed_trigger.jl.txt:108`):

```
antarctic_surface_temperature[t] = (global_surface_temperature[t-1] - ais_temperature_intercept)
                                   / ais_temperature_coefficient
```

so $a \equiv dT_{\text{ant}}/d\text{GMST} = 1 / \text{ais_temperature_coefficient}$. The recalibration samples `amp` and maps it **with the pre-industrial anchor held fixed** (`calibrate_mcmc_ext.jl:274–276`):

```
ais_temperature_coefficient = 1.0 / amp
ais_temperature_intercept   = - AIS_TANT0 / amp      # AIS_TANT0 = -15.42 / 0.8365 = -18.435
```

Decoding the constants (`calibrate_mcmc_ext.jl:164–171, :182`):

- **0.8365** = the original DAIS `ais_temperature_coefficient` (Shaffer 2014 GMST→ T_{ant} regression slope). Its inverse $1/0.8365 = 1.19546$ is the old **equilibrium amplification** that BRICK 2.0 uses.
- **15.42** = the original `ais_temperature_intercept`.
- **AIS_TANT0 = −18.435 °C** = the pre-industrial Antarctic surface-temperature anomaly anchor (Shaffer’s ≈ -18 °C PI), held constant as `amp` varies so that only the *slope*, not the PI baseline, moves.

Justification

The 1.196 slope is an **equilibrium/paleo** relationship; the DAIS fast-dynamics threshold is crossed on a **transient** trajectory. Reading the Antarctic-amplification **secant** (Antarctic $\div$ global cumulative warming since pre-industrial) from **34 CMIP6 models in the land frame** (AIS `tas`), it settles to $\approx 1.06\text{--}1.10$ at the DAIS-crossing-relevant 2.5–3.5 K of global warming. BRICK-AM adopts **$N(1.08, 0.15)$** . The full argument — including why the phase-2 prior’s 0.95 was too low (it came from Xie 2022’s PAI1 trend-ratio computed in an all-cells rather than land frame) — is the authoritative write-up:

`notes/writeup_2026-07-22_a6_amplification_for_tony.{md,pdf}` — the primary justification document (sent to Tony Wong).

Supporting CMIP6 diagnostics: `python/diag_pai_cmip6_time.py`, `diag_pai_deck.py` (idealized 1pctCO2 / abrupt-4×CO2), `diag_pai_ohc.py`, `diag_pai_denominator.py`, `diag_pai_mask_sensitivity.py`.

Code

Location	Role
<code>calibrate_mcmc_ext.jl:274–276</code>	the <code>amp</code> → (coefficient, intercept) mapping with PI anchor preserved
<code>calibrate_mcmc_ext.jl:176–184</code>	the A6 prior — file default <code>N(0.95, 0.10)</code> on <code>[0.70, 1.25]</code>
<code>calibrate_mcmc_ext.jl:59–69</code>	CLI: <code>--amp-mu=</code> , <code>--amp-sigma=</code> , <code>--tag=</code> (note the <code>=</code> -form), and <code>--amp-equilibrium</code>

The **BRICK-AM value** `N(1.08, 0.15)` is applied via CLI override (`--amp-mu=1.08 --amp-sigma=0.15 --tag=extA108`), which also widens the sampling bounds to $\mu \pm 3\sigma = [0.63, 1.53]$ (`:178–181`). The equilibrium sensitivity run pins `a = 1.19546` ($\sigma = 0.002$) via `--amp-equilibrium --tag=extA6eq`.

Flag: the file’s *default* prior is 0.95 (the superseded phase-2 value). BRICK-AM is defined by the CLI override to 1.08 — the tag `extA108` is what makes a posterior “BRICK-AM.”

5. The sub-annual DAIS crossing correction (numerical)

This is **not a model update** — it is a numerical fix, and it matters **only for the pulse marginal**, not for levels or the calibration.

The problem. DAIS disintegration triggers when Antarctic temperature crosses `temperature_threshold`. The stock integrator flips this on **whole-year** boundaries, so a 10 GtCO₂ pulse that advances a draw’s crossing by a *fraction* of a year is rounded to zero for most draws — biasing the finite-difference pulse **median** $\sim 3\text{--}4\times$ low (it collapses to the non-AIS floor) while the mean is only mildly affected.

The fix. Scale the crossing-year disintegration by the fraction of the year spent above threshold, using a backward-mean Antarctic temperature:

```
frac = clamp( (T_sm - threshold) / max(dT_sm, 1e-4), 0, 1 )
disintegration_rate[t] = -λ · 24.78e15 / 57.0 · frac
```

Where it lives. The canonical patched component is saved in-repo at

`julia/patches/antarctic_icesheet_smoothed_trigger.jl.txt` (`:181–200`). It is applied by hand to the loaded depot file

`~/julia/packages/MimiBRICK/edplP/src/components/antarctic_icesheet_component.jl` (the `get_model` path does not read the local repo copy). Apply/restore is **manual** — `chmod u+w` → paste the `frac` block → `run` → `restore from backup` → `chmod u-w` — and every pulse diagnostic **guards and aborts if the patch is absent** (`diag_subannual_pulse_means.jl:13–14`, `diag_decomposition_pulse.jl:13–14`).

It is not used during calibration. The 1850–2026 fit window never crosses the DAIS threshold (ΔT_{glob} only reaches ~ 1.3 K vs the ~ 2.9 K crossing), so the patch changes calibration by nothing and levels by $< 1\%$. It is applied only for the pulse-marginal diagnostics. Recipe: `notes/handoff_2026–07–22_a108_recalibration.md:53–59`.

6. The calibration + deployment pipeline

Environment: `julia --project=julia_v2` (MimiBRICK v2.0.0). Sampler:

`RobustAdaptiveMetropolisSampler.RAM_sample(...; opt_α = 0.234)`. Free parameters: 39 (29 physical + 10 AR(1) nuisance).

```
# 1. Build the reconciled/extended targets (-> recalib_targets_ext.csv + _sources.csv)
python python/prep_recalib_targets_ext.py
```

```
# 2. MCMC: 4 chains x 2,000,000 iters, seeds 2026–2029, over-dispersed starts,  $a \sim N(1.08, 0.15)$ 
for SEED in 2026 2027 2028 2029; do
  julia --project=julia_v2 julia/calibrate_mcmc_ext.jl 2000000 $SEED \
    --overdisperse --amp-mu=1.08 --amp-sigma=0.15 --tag=extA108 &
done; wait
```

```
# 3. Convergence gate on the SLR deliverable (writes outputs/mcmc/slr_convergence_extA108.csv)
julia --project=julia_v2 julia/diag_slr_convergence_by_chain.jl --tag=extA108
```

```
# 4. Accept-on-SLR + write the deployable posterior subsample
julia --project=julia_v2 julia/postprocess_mcmc_ext.jl --tag=extA108 --accept-slr
#   -> data/MimiBRICK/parameters_subsample_brick_mengel_extA108.csv (10,000 rows x 39 cols)
```

```
# 5. Sub-annual pulse projection (REQUIRES the depot patch of $5; aborts otherwise)
julia --project=julia_v2 julia/diag_subannual_pulse_means.jl
```

Production driver: `run_vnext_production.sh`. Equilibrium sensitivity: `run_A6eq_sensitivity.sh` (`--amp-equilibrium`, tag `extA6eq`). The MCMC reuses the phase-2 `outputs/mcmc/adapted_cov_ext.csv` proposal and `overdispersed_starts.csv`.

Priors consumed by step 2: `outputs/param_priors.csv` (22 independent-Gaussian physical priors), `outputs/paleo_geo_prior_ton.csv` (the joint 7-param DAIS-geometry prior + 7×7 correlation, in the identified runoff-onset coordinate), and `outputs/recalib_central_row.csv` (the medoid post-#93 member used as the model base and geometry start point).

7. Validation & results

Convergence — `outputs/mcmc/slr_convergence_extA108.csv` (the SLR deliverable, 4 chains):

Horizon	R	ESS
2100	1.0035	1578
2150	1.0025	1588

Acceptance ≈ 0.238 ; several nuisance AIS-geometry marginals never reach $\hat{R} < 1.05$ (a compensating ridge), but the **projected SLR is converged** — acceptance is gated on the deliverable, not the nuisance marginals.

Level decomposition BRICK 2.0 $\rightarrow$ BRICK-AM (`outputs/decomposition_ssp245.csv`, GMSL cm, SSP2-4.5 median):

horizon	BRICK 2.0	+ Mengel	+ obs/recalib	+ amplification	= BRICK-AM
2100	70.3	-1.1	-3.1	-17.4	48.7
2150	136.2	-4.4	+0.5	-24.0	108.4

The amplification is $\approx 85\%$ of the level change; it moves BRICK-AM from above-AR6 to **mid the AR6 likely range** (0.40–0.60 m @2100).

Pulse decomposition (`outputs/decomposition_pulse.csv`, MAGICC ensemble + sub-annual patch, $\times 10^{-3}$ cm/GtCO₂ @2100): the amplification (–) and recalibration (+) terms partially cancel, so the per-tonne response barely moves (median 21.4 $\rightarrow$ 16.6, mean 21.4 $\rightarrow$ 18.4). The sub-annual patch (§5) is the single largest element of the pulse *median* (+15.8), negligible for the level and the mean. See the cross-model artifact’s decomposition panel.

Projection-file caveat. The regenerated `outputs/proj_ssps_mengel_ext_*` and `postpred_ext_*` files are the **pre-amplification** ext posterior (a $\approx$ equilibrium; SSP2-4.5 @2100 p50 = 75.4 cm). The BRICK-AM (a = 1.08) result lives in the decomposition CSVs and the cross-model artifact (@2100 median 48.7 cm), **not** in a regenerated SSP/postpred file. Regenerate `posterior_predictive_ext.jl` with `--tag=extA108` if a fresh BRICK-AM postpred is needed.

Appendix A — File index

Data / inputs

Path	Contents
outputs/recalib_targets_ext.csv (+ _sources.csv)	calibration targets (5 components + total), 1900–2026, cm rel. 1995–2005
data/observations/raw/	GRACE-FO AIS/GIS, GlaMBIE glaciers, NOAA thermosteric, Frederikse 5000-member .nc
data/observations/dangendorf2024_gmsl_annual.csv, nasa_gmsl_annual.csv	total-GMSL obs (Dangendorf + NOAA STAR)
data/observations/fair_mean_gmst_ssp245harm.csv, ..._ohc_ssp245harm.csv	FaIR forcing (calib1.4.5, harmonized SSP2-4.5)
outputs/param_priors.csv, outputs/paleo_geo_prior_ton.csv, outputs/recalib_central_row.csv	priors + medoid base
data/MimiBRICK/parameters_subsample_brick_mengel_extA108.csv	the deployable BRICK-AM posterior (10k × 39)

Code

Path	Role
python/prep_recalib_targets_ext.py	build targets + σ
julia/glaciers_mengel_component.jl	Mengel glacier component
julia/brick_mengel.jl, julia/brick_param_updates.jl	model assembly + param mapping
julia/calibrate_mcmc_ext.jl	likelihood, priors, amplification mapping, MCMC
julia/postprocess_mcmc_ext.jl, julia/diag_slr_convergence_by_chain.jl	accept-on-SLR + convergence
julia/patches/antarctic_icesheet_smoothed_trigger.jl.txt	canonical sub-annual DAIS patch
julia/diag_decomposition.jl, julia/diag_decomposition_pulse.jl, julia/diag_subannual_pulse_means.jl	decomposition + pulse diagnostics
run_vnext_production.sh, run_A6eq_sensitivity.sh	production drivers
FaIRtoFrEDI/build_fair_mean_v145.py	forcing build

Appendix B — Notes & write-ups (all in `SLR-RFF-BRICK/notes/`)

File	Covers
<code>writeup_2026-07-22_a6_amplification_for_tony.{md,pdf}</code>	the amplification justification (primary)
<code>writeup_2026-07-21_brick_fm_vs_wong_brick.md</code>	BRICK-FM vs Wong BRICK comparison
<code>handoff_2026-07-22_a108_recalibration.md</code>	the $a=1.08$ cold-start recipe + patch procedure
<code>handoff_2026-07-20_phase2_calibration_wired.md</code> , <code>prerun_summary_2026-07-20_phase2_calibration.md</code>	phase-2 (A2/A4/A5/A6) wiring
<code>handoff_2026-07-19_brick_fm_improvement_roadmap.md</code> , <code>handoff_2026-07-18_brick_mengel_vnext.md</code>	roadmap + v-next kickoff
<code>handoff_2026-06-12_brick_mengel_calibration.md</code> , <code>handoff_2026-06-13_brick_mengel_post2018_extension.md</code> , <code>handoff_2026-06-13_brick_mengel_postpred_projections.md</code>	early Mengel provenance

Appendix C — Provenance commits (branch `brick-mengel-vnext`)

`7423ab3` port Mengel · `841e679` LIA offset · `047b27a` 2τ + LWS lock · `180aa31` acquire extension data + splice · `636f849` ext-target calibration · `9dd7d95` free DAIS geometry · `a954701` phase-2 A2/A4/A5/A6 · `cda7ca2` M3 Dangendorf+STAR + ensemble σ · `2b43e35` `--amp-equilibrium` · `ee0685b` wire $a=1.08$ · `69a2bfb` pulse decomposition (HEAD).